

For this assignment the following MATLAB code will be required:

<http://www.mdunlop.org/cms107/assignment1.zip>

An Introduction to Clustering

A common problem in unsupervised learning is that of clustering: grouping a collection of data or objects into clusters, such that objects within a cluster are more similar to one another than they are to objects outside of the cluster. Typical applications include image segmentation, anomaly detection and social network analysis. Figure 1 illustrates the idea of clustering in the case of planar data: 800 points are given distributed in the unit square, and a clustering algorithm attempts to classify these points into four clusters based on the distance between them. This is a simple example for illustration – in general the data may be very high dimensional, and there may be nonlinear structures underlying the data that need to be inferred.

One approach to clustering is to find a partition of the data into k clusters such that the within-cluster variance is minimal, i.e. to find a minimizer of the functional

$$I(S_1, \dots, S_k) = \sum_{j=1}^k \sum_{x \in S_j} \|x - \mu_j\|_2^2$$

which is defined on partitions of the data. Here μ_j is the mean of the points in S_j . This method is known as *k-means clustering*, and whilst computationally expensive to solve exactly, there are a number of heuristic algorithms that make the computation feasible [1].

We focus on *spectral clustering*, which involves looking at the eigenvectors and eigenvalues of a graph Laplacian associated with the data. It is outlined in the next section, and described in more detail in the tutorial [2]. In this assignment we first consider clustering for synthetic data, so that there are ‘true’ clusters that are to be determined, generated from a known statistical model. We then look at an example with real data: classifying political party affiliation from voting records.

Spectral Clustering

We are given data $x_1, \dots, x_n \in \mathbb{R}^d$. We can think of these data points as the vertices of an undirected graph, which we will call the *similarity graph* and denote¹ G . Assume that the edge between vertices x_i and x_j has weight $w_{ij} \geq 0$, where w_{ij} represents the similarity of the points x_i and x_j . A typical way to define these weights is by

$$w_{ij} = \exp\left(-\frac{d(x_i, x_j)^2}{\ell^2}\right) \quad (1)$$

¹We will denote by G both the entire graph and just its vertices, as this should not lead to any ambiguity in the contexts considered here.




Figure 1: An example of clustered data in the plane (left) and an attempt to partition these data into four clusters (right).

for some metric d on \mathbb{R}^d and scalar $\ell > 0$. The parameter ℓ is a length-scale parameter related to the size and separation distance between clusters; note that the smaller the value of ℓ , the greater the sensitivity of w_{ij} to the distance between the points x_i and x_j . In particular, note that in the limit $\ell \rightarrow 0$ the weights become one or zero depending on whether the two points being considered are coincident or not.

Define also the degree δ_i of the vertex x_i by

$$\delta_i = \sum_{j=1}^n w_{ij}.$$

We form two matrices $W, D \in \mathbb{R}^{n \times n}$ using the weights and the degrees:

$$W_{ij} = w_{ij}, \quad D_{ij} = \begin{cases} \delta_i & i = j \\ 0 & i \neq j \end{cases}.$$

Using these matrices we can form three further matrices, called graph Laplacians:

Definition 1 (Graph Laplacians). *Let $D, W \in \mathbb{R}^{n \times n}$ be the degree and weight matrices above. Define the following three graph Laplacian matrices:*

- $L := D - W$
- $L_{\text{sym}} := D^{-1/2} L D^{-1/2} = I - D^{-1/2} W D^{-1/2}$
- $L_{\text{rw}} := D^{-1} L = I - D^{-1} W$

L is called the unnormalized graph Laplacian, L_{sym} is called the symmetric normalized graph Laplacian and L_{rw} is called the random walk normalized graph Laplacian.

The eigenvalues and eigenvectors of these Laplacians turn out to be a very useful tool for the identification of clusters, as will be seen in the problems that follow.

Problem 1. Theory (40 points)

We focus here on the unnormalized Laplacian L ; similar results hold for L_{sym} and L_{rw} . In what follows we assume nothing about the form of the weights w_{ij} except that they are symmetric and non-negative. Additionally, given a subset of vertices $A \subseteq G$ we define $\mathbf{1}_A \in \mathbb{R}^n$, the indicator vector of A , by

$$(\mathbf{1}_A)_j = \begin{cases} 1 & x_j \in A \\ 0 & x_j \notin A \end{cases}.$$

(a) Show that

(i) for any $f \in \mathbb{R}^n$,

$$\langle f, Lf \rangle = \frac{1}{2} \sum_{i,j=1}^n w_{ij}(f_i - f_j)^2;$$

(ii) L is symmetric and positive semi-definite; and

(iii) the smallest eigenvalue of L is 0, with corresponding constant eigenvector $\mathbf{1}_G$.

Deduce that L has n non-negative real eigenvalues $0 = \lambda_1 \leq \lambda_2 \leq \dots \leq \lambda_n$.

(b) Given two vertices $x_i, x_j \in G$, we will say that there exists a path between x_i and x_j if there is a set of vertices $x_{m_1}, \dots, x_{m_p} \in G$ such that the weights

$$w_{im_1}, w_{m_1m_2}, w_{m_2m_3}, \dots, w_{m_{p-1}m_p}, w_{m_px_j}$$

are all strictly positive. We will say that $A \subseteq G$ is a connected component of G if both $x_i, x_j \in A$ implies that there exists a path between x_i and x_j , and if only one of $x_i, x_j \in A$ implies that there does not exist a path between x_i and x_j . We will say that G has k connected components if there is a disjoint collection of connected components $A_1, \dots, A_k \subseteq G$ whose union is G .

(i) Consider Figure 2, representing a graph with 5 vertices. Which of the following are connected components?

- $\{x_1, x_4, x_5\}$
- $\{x_1, x_2\}$
- $\{x_1, x_2, x_3\}$

How many connected components does the graph have?

(ii) Let $A \in \mathbb{C}^{(n_1+n_2) \times (n_1+n_2)}$ have block diagonal form

$$A = \begin{pmatrix} A_1 & 0 \\ 0 & A_2 \end{pmatrix}$$

for $A_1 \in \mathbb{C}^{n_1 \times n_1}$ and $A_2 \in \mathbb{C}^{n_2 \times n_2}$. Assume that A_1 has eigenvalues $\lambda_1^1, \dots, \lambda_{n_1}^1$ with corresponding eigenvectors $v_1^1, \dots, v_{n_1}^1 \in \mathbb{C}^{n_1}$, and A_2 has eigenvalues $\lambda_1^2, \dots, \lambda_{n_2}^2$ with corresponding eigenvectors $v_1^2, \dots, v_{n_2}^2 \in \mathbb{C}^{n_2}$. What are the eigenvalues and eigenvectors of A ? How does this generalize to a larger number of blocks?

(iii) Show that the 0 eigenvalue of L has geometric multiplicity 1 if and only if the graph G has one connected component.




Figure 2: An example of a graph with weighted edges. A dashed line between x_i and x_j corresponds to $w_{ij} = 0$; a solid line corresponds to $w_{ij} > 0$.

- (iv) Show that the 0 eigenvalue of L has geometric multiplicity k if and only if the graph G has k connected components $\{A_m\}_{m=1}^k$. Show that in this case the eigenspace of the 0 eigenvalue is spanned by the indicator vectors $\mathbf{1}_{A_1}, \dots, \mathbf{1}_{A_k}$.

Part (b) of the exercise above tells us that when the similarity graph G has k connected components, these components can be perfectly identified given knowledge of the first k eigenvectors of the graph Laplacian L . With a good choice of weights w_{ij} , these connected components in G should correspond to the clusters in \mathbb{R}^d . Choice of such weights is not trivial however, as will be seen in the following problems. Nonetheless, even with imperfect weights this provides useful heuristics for determining the clusters.

Problem 2. Synthetic Data (40 points)

- (a) Implement the three graph Laplacians L , L_{sym} and L_{rw} in MATLAB. Specifically, given an array of data $X \in \mathbb{R}^{n \times d}$, with each row of X representing a data point $x_i \in \mathbb{R}^d$, define a weight matrix $W = \{w_{ij}\}_{i,j=1}^n$ using the form (1) for the weights with the choice of Euclidean distance

$$d(x, y) = \|x - y\|_2,$$

and any length-scale parameter $\ell > 0$. Then define the corresponding degree matrix D , and return the three graph Laplacians as given in Definition 1.

- (b) Load `data2.mat`. The data sets X1, X2, X3 arise from a number of points distributed into a number of clusters in the plane, embedded in \mathbb{R}^{100} and then perturbed by noise in all directions.
- (i) Using the `eig` function, compute the eigenvectors of the three Laplacians for a variety of choices of $\ell > 0$, using the data X1. Problem 1(b) told us that information about clusters should lie in the first few eigenvectors – by looking at these, attempt to infer the number of clusters and their elements. How sensitive is the clustering to the choice of ℓ , and the choice of Laplacian?






Figure 3: A selection of images of the digits 1, 4, 9 from the MNIST data set to be classified in Problem 3.

- (ii) Instead of a fixed length-scale parameter $\ell > 0$, the length scale can be inferred from the data. Fix $K \in \mathbb{N}$. The self-tuning weights of [3] are defined by

$$w_{ij} = \exp \left(-\frac{\|x_i - x_j\|_2^2}{\ell_i \ell_j} \right) \quad (2)$$

where $\ell_j = \|x_j - x_K\|_2$, and x_K is the K th nearest neighbor of x_j . An example MATLAB implementation of these weights is provided in `weights_st.m`. Using this, or your own implementation if you wish, repeat the clustering from part (i) using the self-tuning weights with $K = 10$.

- (iii) The three sets of data X_1 , X_2 , X_3 arise from the same clusters, perturbed by increasing levels of noise. How does the noise level affect the ability to determine the clusters?
- (iv) Illustrate the output of some of the above clustering via a projection of the clustered data onto its first two components.

Problem 3. Real Data (20 points)

- (a) Load `data3.mat`. The matrix X contains 3085 rows, each corresponding to a grayscale image of a handwritten digit 1, 4 or 9. The images are reshaped from 20×20 matrices to vectors in \mathbb{R}^{400} . A selection of these images are shown in Figure 3. The aim of this problem is to separate the images into three clusters corresponding to these digits.
 - (i) Choose one of either L , L_{sym} or L_{rw} , and let the weights take the form (2). Look at the eigenvectors of the Laplacian for a number of different choices of the parameter K in the weights – does the choice of K significantly influence the quality of the clustering?
 - (ii) Can all three clusters be identified? Is one cluster easier to identify than the other two, and if so why do you think this is?

References

- [1] Anil K Jain. Data clustering: 50 years beyond k-means. *Pattern recognition letters*, 31(8):651–666, 2010.
- [2] Ulrike von Luxburg. A tutorial on spectral clustering. *Statistics and Computing*, 17(4):395–416, 2007.
- [3] Lihi Zelnik-Manor and Pietro Perona. Self-tuning spectral clustering. In *Advances in neural information processing systems*, pages 1601–1608, 2005.